



I'll Be Backoff: Benchmarking AI-Powered Platform Engineering Tools

Annie Talvasto,
Sr. Manager, Product
Marketing & Community
Upbound

AI & Platform Engineering



Why should you
care?



What are the
options?



Landscape: 4 Major Categories

1. AI SRE (Incident Intelligence)

Tools that reason, triage, investigate, and suggest or coordinate fixes.

2. AI Platform Engineering / SDLC Automation

Tools that automate workflows, tickets, build-test-release, scorecards, and developer experience.

3. Infrastructure Automation (IaC & Control Planes)

AI for infra mutation and enforcement.

4. Agent Runtimes (Agentic Infra)

Where agents actually run, the execution substrate for agentic workloads.

AI SRE (Incident Intelligence)



Datadog Bits AI SRE

- Fully autonomous investigation inside Datadog
- Telemetry-native: logs, metrics, traces, Watchdog
- Incident summaries, coordination, postmortems
- Best for orgs already on Datadog

Phoebe

- High-accuracy log/trace/query reasoning
- Predicts incidents and generates PR fixes
- Builds knowledge graphs for faster debugging
- Vendor-agnostic read-only ingestion

Cleric

- Transparent hypothesis-driven investigations
- Shows reasoning chains and validated evidence
- Engineer collaboration for fixes
- Read-only by design


Resolve AI

- Autonomous incident investigation + RCA
- Cross-system evidence collection
- Fix proposals + human approval loop
- Multi-agent orchestration

Traversal

- Evidence-backed RCA, alert ranking, fast diagnosis
- Designed for on-prem / high-security teams
- Fix suggestions + postmortem automation
- Not focused on SDLC automation

AI Platform
Engineering / SDLC
Automation



Port

- Developer portal with context lake + scorecards
- AI agents automate tickets, actions, and workflows
- Strong SDLC catalog visibility
- Not an incident/RCA tool


SRE AI

- Broad AI across testing, CI/CD, release, security
- Ticket/PR automation, risk detection
- Great for hybrid/remote engineering orgs
- Moderate depth in RCA

Harness AI SRE

- Incident triage + RCA via change intelligence
- Runbook automation across the Harness platform
- Strong SDLC-wide visibility
- Best value if using Harness CI/CD

Infrastructure Automation (IaC & Control Planes)



Pulumi Neo

- AI agent for infrastructure-as-code changes
- Generates PRs, upgrades, fixes, compliance updates
- Strong governance via previews + policies
- Deep Pulumi ecosystem integration

Upbound

- Declarative, self-healing infra via Crossplane 2.0
- Deterministic APIs for humans + AI agents
- Policy enforcement + drift correction
- Foundation layer; not an SRE/RCA system
- Build on top of the trusted reconciliation technology/controllers as Kubernetes

Azure SRE Agent

- Azure-native autonomous remediation agent
- Always-on monitoring + active incident fixes
- Works only on Azure resources
- Infra-level autonomy, limited RCA reasoning

Agent Runtimes (Agentic Infra)



Kagent

- MCP/ADK agents running inside Kubernetes
- Cluster-native workflows, tools, tracing
- Extensible via CRDs + open standards
- Not an SRE/RCA product, it's the runtime

Tool	Primary Domain	What It Actually Does	Best For
Datadog Bits AI SRE	Native observability-driven SRE	Fully autonomous investigations, incident coordination, summaries, postmortems	Organizations already on Datadog
Phoebe	Proactive + reactive AI troubleshooting	Predict issues, RCA, PR fixes, high-accuracy log/search models	High-scale orgs with rich observability pipelines
Cleric	Explainable AI SRE	Hypothesis-driven RCA, structured reasoning, human-collab debugging	Teams needing transparent, explainable agents
Resolve AI	AI SRE (multi-agent)	Incident investigation, RCA, code/infra context, fix proposals	SRE teams needing deep incident reasoning
Traversal	AI SRE with privacy focus	Alert ranking, RCA, remediation suggestions, postmortems, on-prem/bare-metal friendly	Enterprise SRE teams with compliance constraints
Port	AI-enhanced IDP + workflow automation	SDLC catalog, scorecards, ticket automation, custom AI agents	Large engineering orgs with complex workflows
SRE.ai	Broad engineering automation	Testing, CI/CD insights, SDLC orchestration, predictive workflows	Hybrid/remote teams needing lifecycle automation
Harness AI SRE	Incident mgmt + SDLC change correlation	Triaging, RCA via change intelligence, runbook automation	Companies using Harness CI/CD
Pulumi Neo	AI for IaC automation	Infra code generation, upgrades, compliance fixes, PR automation	Platform engineers using Pulumi
Upbound	Intelligent control plane (foundation)	Declarative infra APIs, drift correction, policy enforcement	Org building AI-native infra & CPaaS
Azure SRE Agent	Cloud-native auto-mitigation	Azure resource monitoring + automated fixes via AAUs	Teams fully on Azure needing auto-reliability
kagent	Kubernetes-native agent runtime	Run MCP/ADK agents <i>inside</i> clusters; agent workflows; observability	Platform teams wanting open standards for agentic infra

Tool	Grounding Source	Strength
Bits AI (Datadog)	Full Datadog telemetry graph	Very strong telemetry-native grounding
Phoebe	Proprietary cross-silo log/trace/search models	High-accuracy reasoning across systems
Cleric	Logs, metrics, code changes	Explainable evidence-based grounding
Resolve AI	Logs, metrics, traces, code, infra maps	Broad multi-source grounding
Traversal	Telemetry + deploy history	Precise incident signal grounding
Port	SDLC metadata, org catalog, workflows	Strong organizational context
SRE.ai	CI, code, tests, SDLC process data	Broad SDLC context
Harness AI SRE	CI/CD + change context	Good for SDLC-to-incident correlation
Pulumi Neo	IaC repos, Pulumi state, policies	Strong code-level grounding
Upbound	Control plane state, CRDs, policies	Deterministic infra grounding
Azure SRE Agent	Azure Monitor, resource health	Strong Azure-native grounding
kagent	K8s resources, MCP tools, ADK	Strong cluster-native grounding

Tool	Autonomy Level	What It Can Autonomously Do
Bits (Datadog)	High for RCA, low for changes	Autonomous reasoning + incident coordination
Phoebe	High	Predictive fixes, PR generation
Cleric	Medium	RCA + explanations, not infra writes
Resolve AI	Medium-high	Multi-agent RCA + fix proposals
Traversal	Medium-high	RCA + remediation suggestions
Port	Medium	Execute actions, workflows, route tickets
SRE.ai	Medium	Testing, release gating, risk detection
Harness AI SRE	Medium	Execute runbooks + correlate changes
Pulumi Neo	Medium	Generate PRs, run previews, propose changes
Upbound	High (declarative)	Continuous reconciliation + drift correction
Azure SRE Agent	High (infra only)	Auto-mitigations in Azure
kagent	Varies by CRD	Agents execute workflows <i>in-cluster</i>

Tool	Strength Highlights	Limitations
Bits AI	Deepest telemetry-native grounding	Datadog-only ecosystem
Phoebe	Strong proactive detection + accurate PR fixes	Not an incident coordinator
Cleric	Explainable reasoning chain	No automated infra writes
Resolve AI	Multi-agent RCA with code/infra intelligence	No infra drift control; not deterministic
Traversal	Fast RCA + enterprise privacy posture	Not multi-cloud infra automation
Port	Strong workflow engine + context lake	Not RCA/incident-focused
SRE.ai	Full SDLC-level automation + testing intelligence	Wide scope but shallow RCA depth
Harness AI SRE	Change intelligence inside CI/CD	Requires Harness stack
Pulumi Neo	Tight IaC grounding, safe PR automation	Not an SRE system; no observability reasoning
Upbound	Declarative control planes; infra correctness	No RCA, no incident workflows
Azure SRE Agent	Auto-remediation for Azure workloads	Azure-only; limited RCA depth
kagent	Open-standards agent runtime for K8s	No RCA or observability; needs K8s maturity

Landscape: 4 Major Categories

1. AI SRE (Incident Intelligence)

Tools that reason, triage, investigate, and suggest or coordinate fixes.

Examples: Datadog Bits, Phoebe, Cleric, Resolve AI, Traversal

2. AI Platform Engineering / SDLC Automation

Tools that automate workflows, tickets, build-test-release, scorecards, and developer experience.

Examples: Port, SRE.ai, Harness AI SRE.

3. Infrastructure Automation (IaC & Control Planes)

AI for infra mutation and enforcement.

Examples: Pulumi Neo, Upbound, Azure SRE Agent.

4. Agent Runtimes (Agentic Infra)

Where agents actually run, the execution substrate for agentic workloads.

Example: kagent

Resources

- Pick what makes most sense for you!
- Intelligent Control Plane
 - <https://www.upbound.io/intelligent-control-plane>
- SREday sessions
 - **AI & SRE: Embracing the Ultimate Frontier of Automation**
<https://youtu.be/7VXZBLROglM?si=2ezbKagZTsR8b3C5>
 - **The Rise of SRE AI Agents**
<https://youtu.be/xWNI4WLxXcw?si=tJkJsOvxEoHsAckY>
 - And many more!



Thank you!

