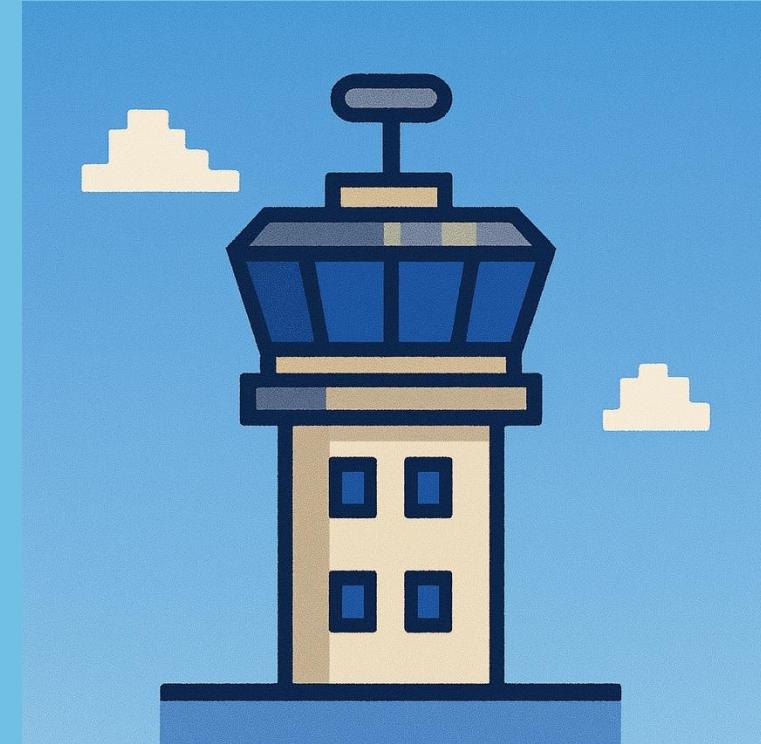




# LLMs on Autopilot: Running AI Agents on Kubernetes With Open Source Tools

Annie Talvasto  
Sr. Manager,  
Product Marketing  
& Community  
Upbound

# Infrastructure & Development in the Age of AI



## What are AI agents?

- Agents are AI systems that don't just assist humans reactively. They plan, decide, and act autonomously toward goals.
- Agents can analyze a request, determine the need for a database, and (attempt) to provision it without human intervention.
- To function effectively, agents need structured, unified access to all operational inputs: documentation, policies, infrastructure APIs.

# Agent Examples

## SRE Agent

Monitors system health, detects anomalies, and automatically remediates incidents to maintain reliability and uptime.

## Provisioning Agent

Automates infrastructure provisioning and configuration based on high-level goals or requests.

## Compliance Agent

Enforces security policies and compliance requirements across cloud workloads

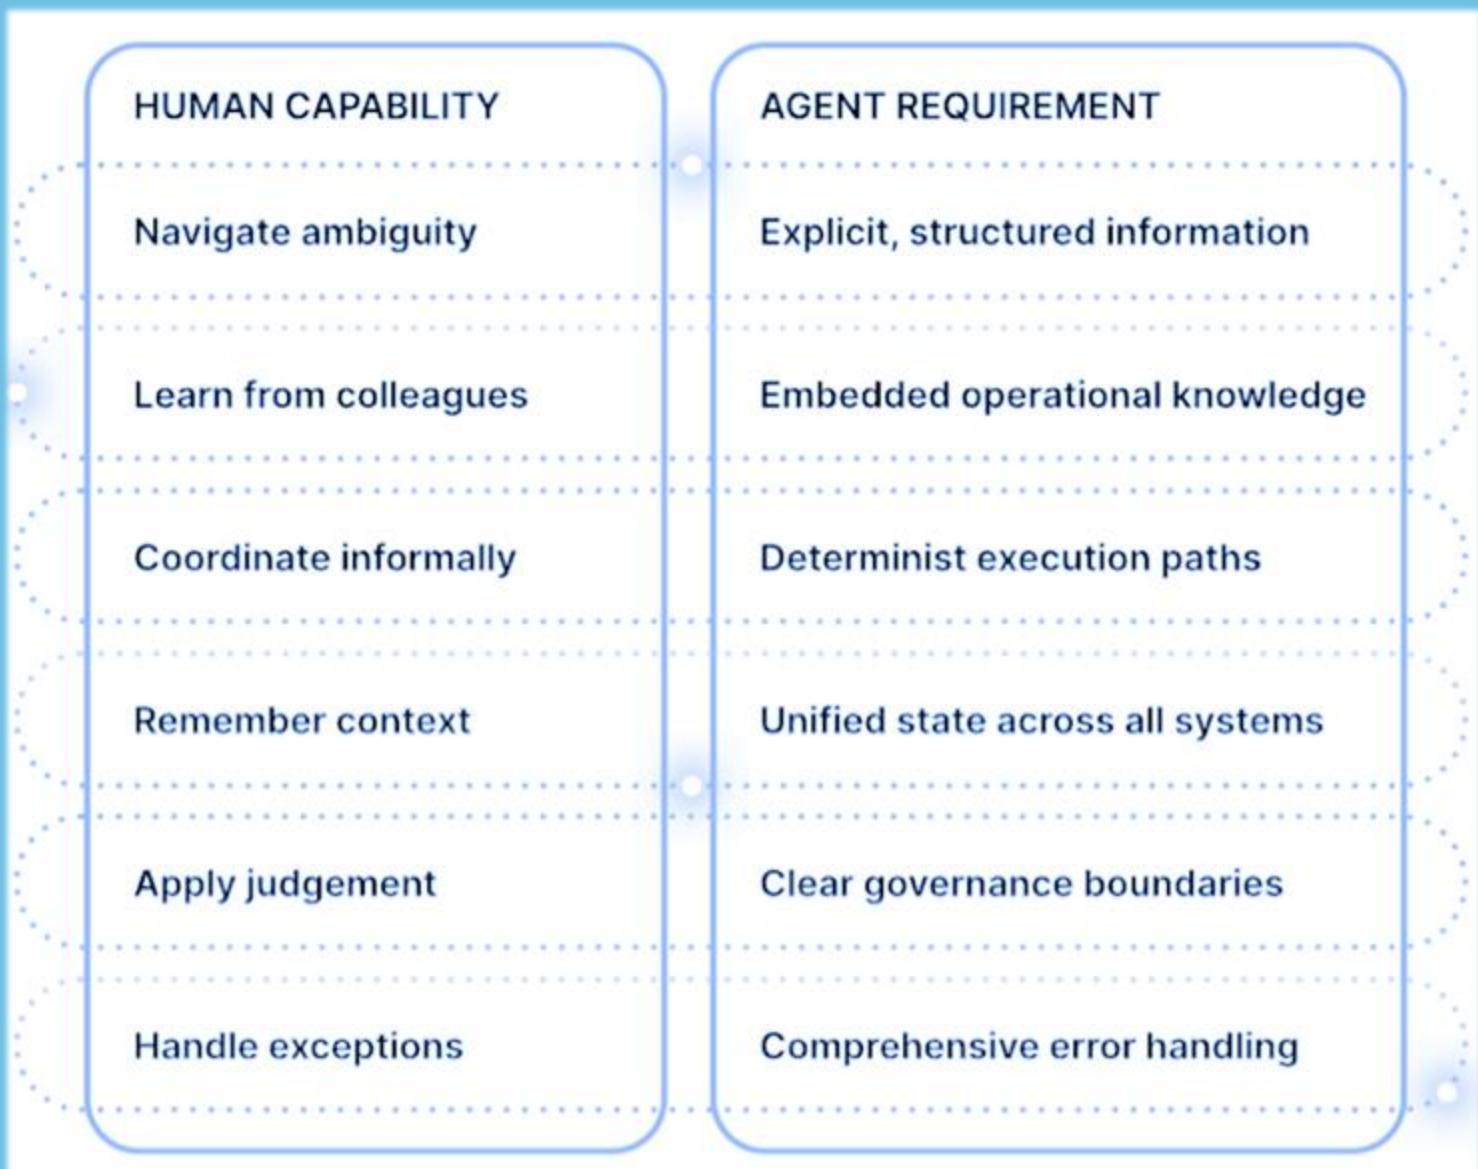
## Cost optimization Agent

Continuously analyzes resource usage and recommends or implements optimizations to reduce cloud spend.

# Agent Requirements Are Fundamentally Different From Humans



# Human capabilities vs agent requirements



**When Terraform partially fails:**

 **Human:** Checks Slack, finds someone mentioned this yesterday, applies workaround

 **Agent:** Sees error state, cannot determine if retry is safe, stops or corrupts state further

**When configuration is incomplete:**

 **Human:** Recognizes what's missing from experience, knows who to ask

 **Agent:** Only knows what's in template, and cannot fill gaps

**When compliance is special:**

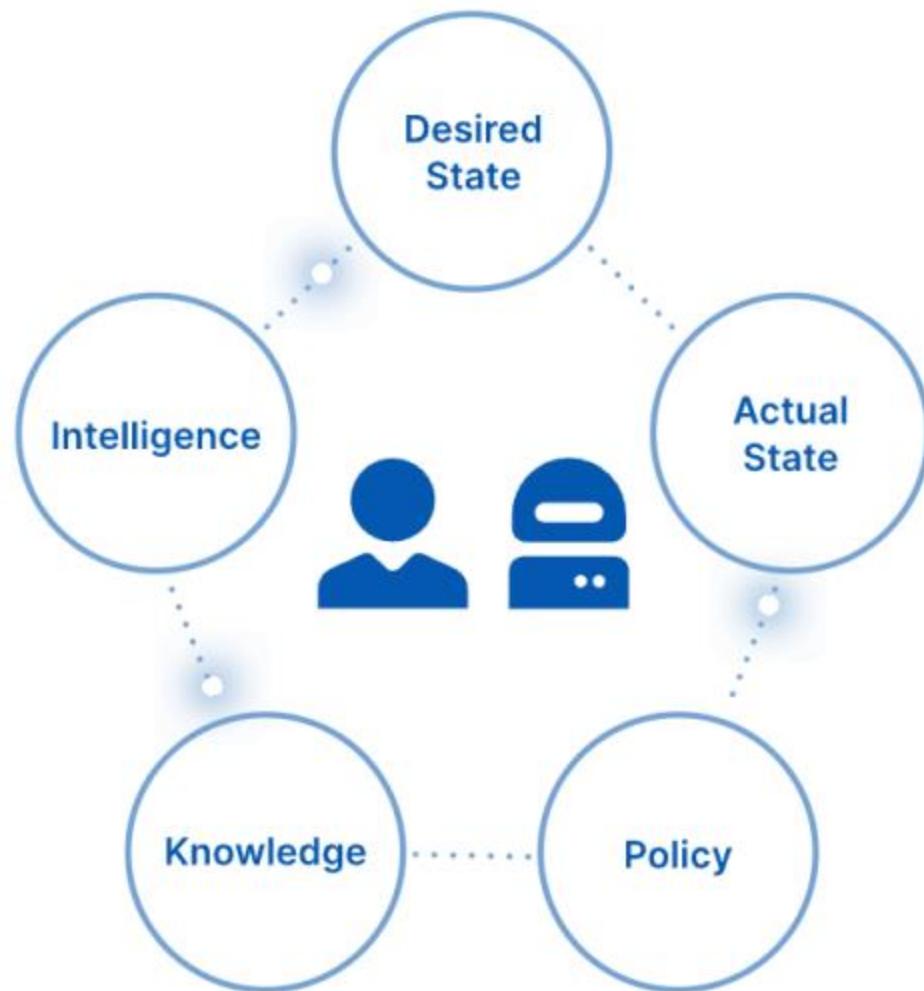
 **Human:** Remembers from training that payment services need extra controls

 **Agent:** Has no knowledge outside automated path

## Why Platforms Break for Agents: Platforms are Built for Humans, Not AI

- **Fragmentation of Operational Elements:** Desired state in Git, actual state in cloud consoles, policies in pipelines, and tribal knowledge in people's heads. Humans bridge this gap; agents cannot.
- **Humans as Error Handlers:** Automation relies on humans to fix brittle paths (rerun scripts, tweak configs). Agents lack judgment and informal workarounds, they either blindly retry or fail.
- **Day-2 Operations & Context:** Scaling, incident response, maintenance depend on institutional memory and ad-hoc coordination. Agents lack context and informal networks, risking chaos in production.

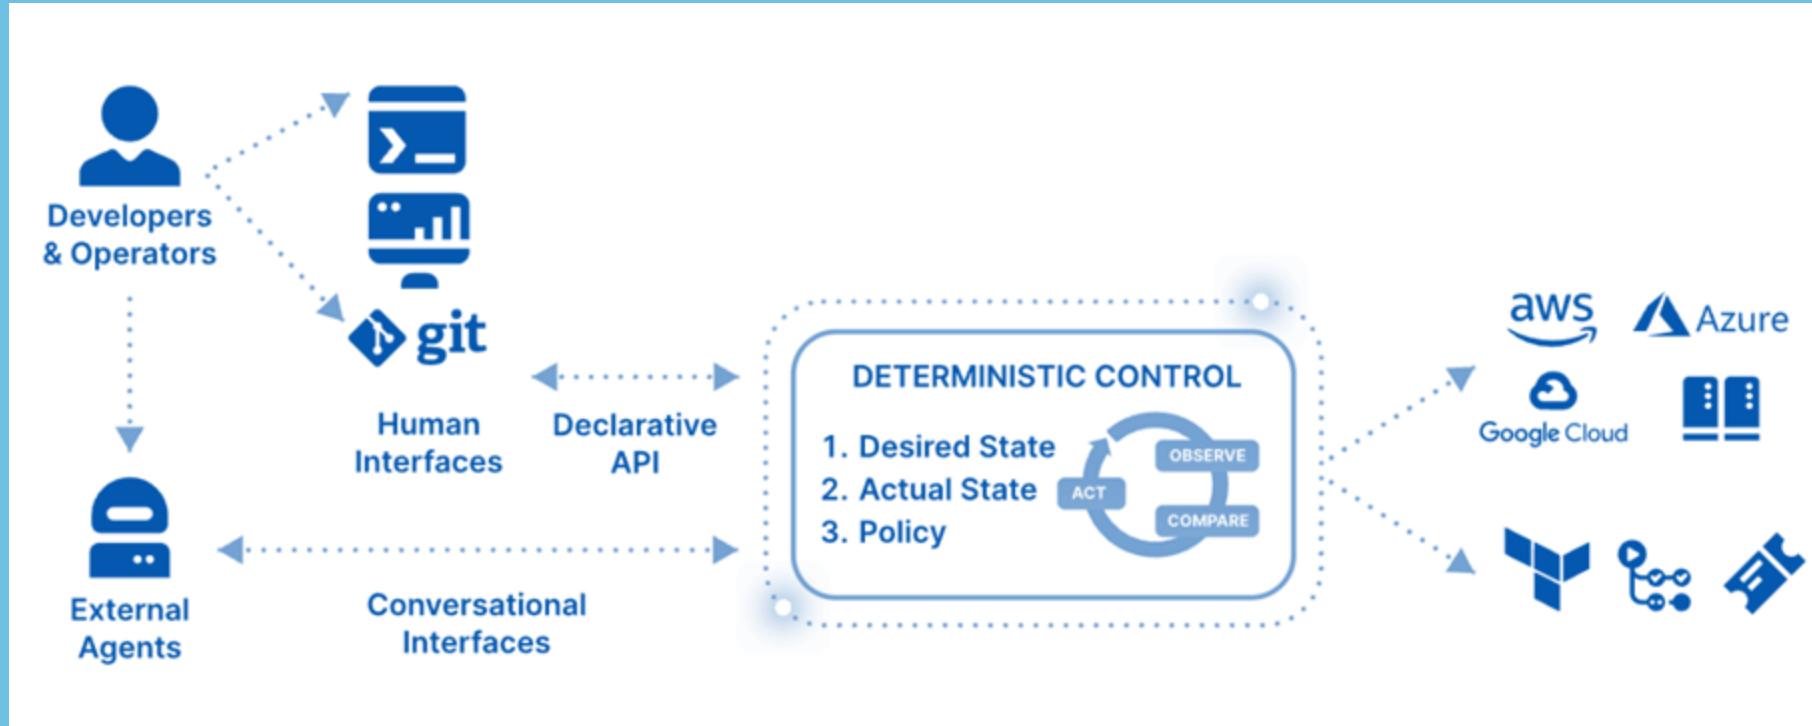
## What is Needed for Agentic ops?



Tools To Use Today and in  
the Future

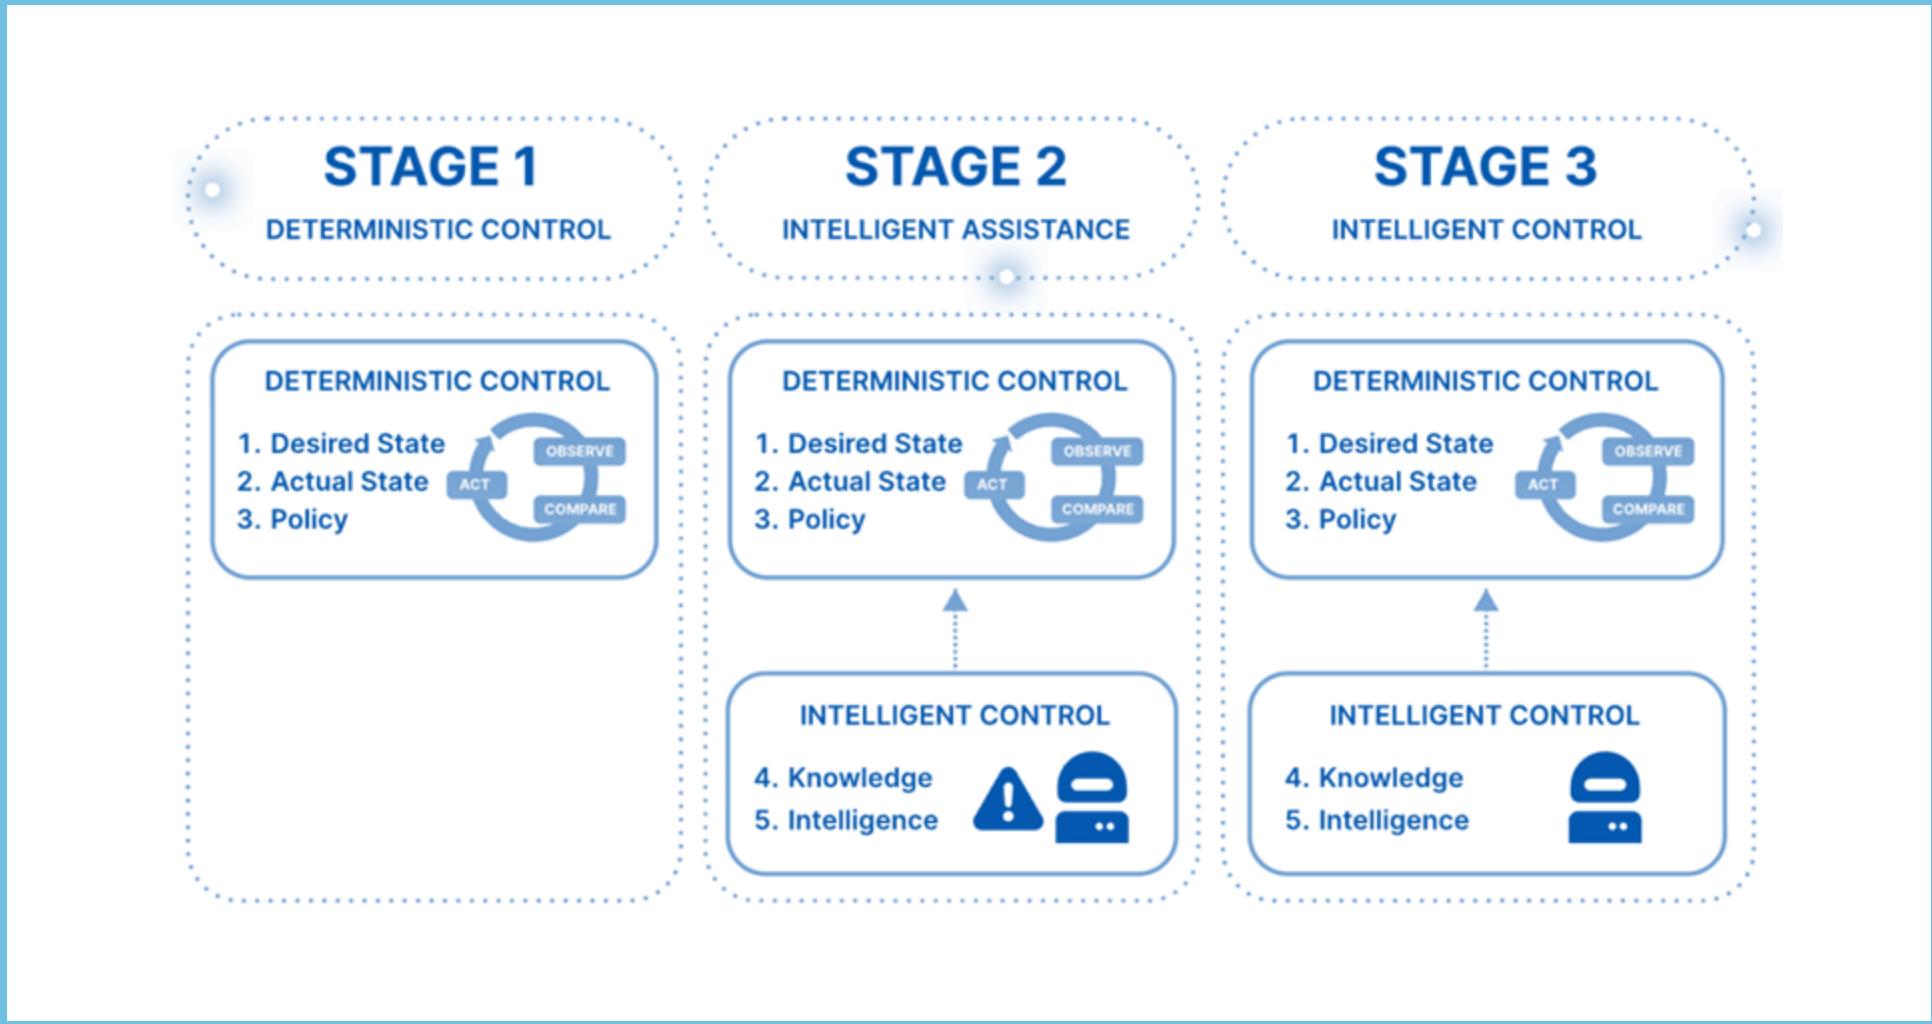


# Infrastructure Orchestration with Crossplane 2.0



- Declarative Desired State
- Composition & Composition Functions
- Continuous Reconciliation & Drift Correction

# Control Plane as Authoritative System of Record



## Prometheus: Monitoring Matters

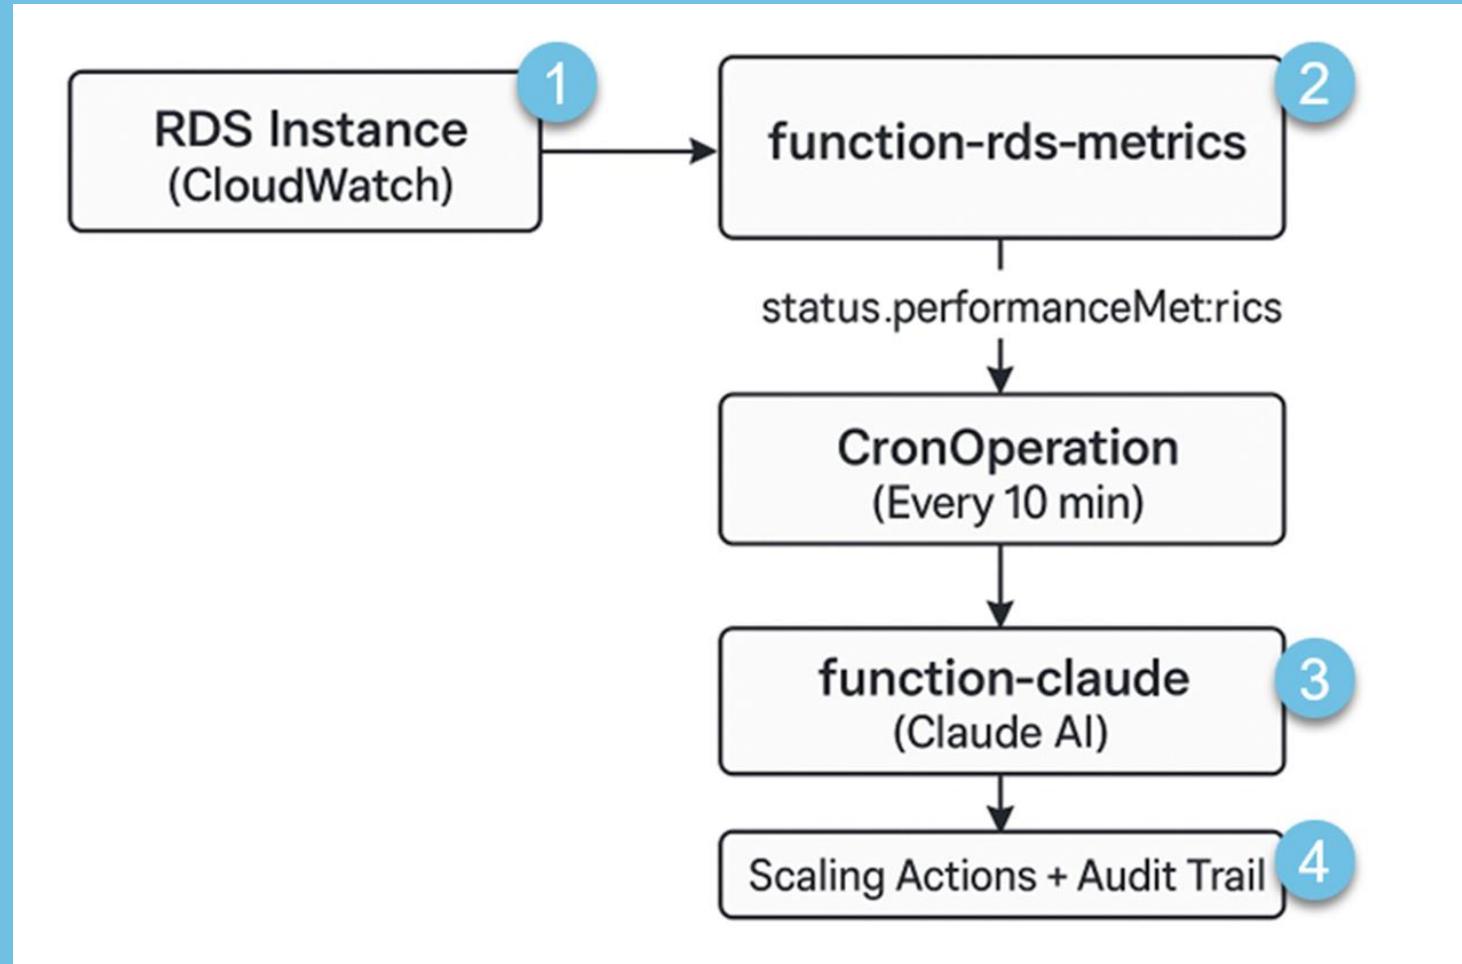
- Agents depend on continuous, reliable metrics
- Metrics → context → decisions
- Prometheus provides the perceptual layer for agent loops
- Observability becomes part of the agent's world model

Demo

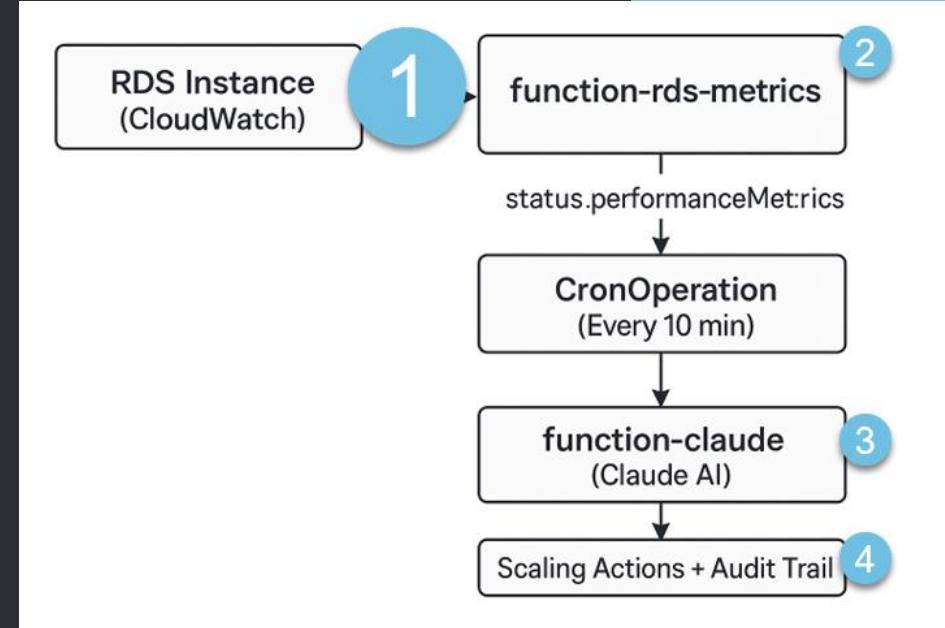


# Demo: An Autonomous Infrastructure Agent on Kubernetes

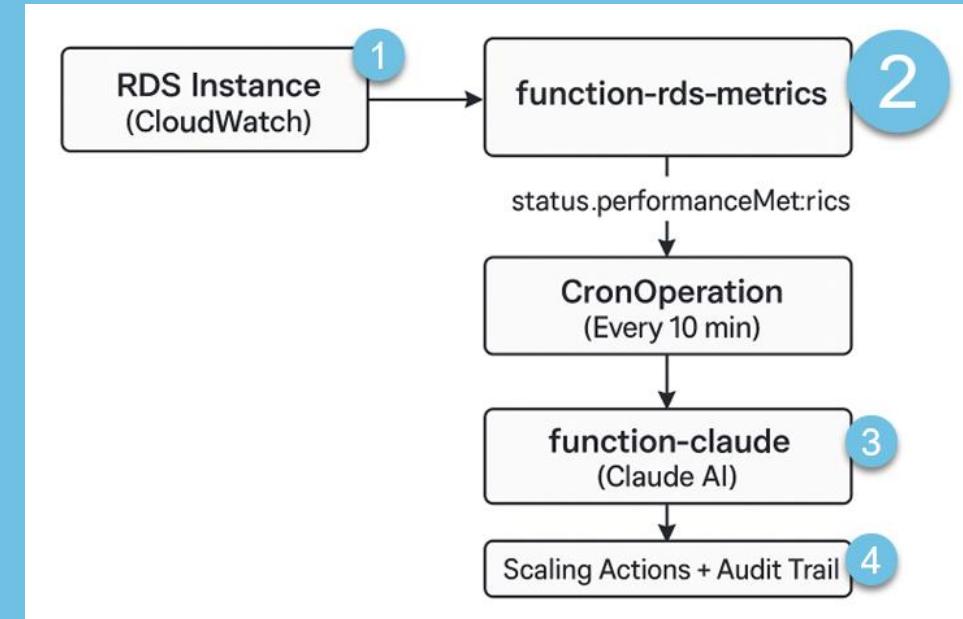
- A real AI agent:
  - Reads metrics
  - Analyzes them with Claude
  - Makes scaling decisions
  - Updates desired state
  - Crossplane enforces it



```
apiVersion: aws.platform.upbound.io/v1alpha1
kind: XSQLInstance
metadata:
  name: rds-metrics-database-ai-scale
  namespace: default
  labels:
    scale: me
spec:
  compositionSelector:
    matchLabels:
      type: rds-metrics
  parameters:
    region: us-west-2
    engine: mariadb
    engineVersion: "10.11"
    storageGB: 5
    instanceClass: db.t3.micro
    autoGeneratePassword: true
    publiclyAccessible: true # Temporarily for stress testing
    passwordSecretRef:
      namespace: default
      name: mariadbsecret
      key: password
    networkRef:
      id: rds-metrics-database-ai-scale
      writeConnectionSecretToRef:
```



```
performanceMetrics:  
  databaseName: rds-metrics-database-ai-scale  
  metrics:  
    CPUUtilization:  
      timestamp: "2025-09-15T12:52:00Z"  
      unit: Percent  
      value: 2.6191111704257306  
    DatabaseConnections:  
      timestamp: "2025-09-15T12:52:00Z"  
      unit: Count  
      value: 0  
    FreeStorageSpace:  
      timestamp: "2025-09-15T12:52:00Z"  
      unit: Bytes  
      value: 2775519232  
    FreeableMemory:  
      timestamp: "2025-09-15T12:52:00Z"  
      unit: Bytes  
      value: 188219392  
    ReadIOPS:  
      timestamp: "2025-09-15T12:52:00Z"  
      unit: Count/Second  
      value: 0.01667722891164404  
    WriteIOPS:  
      timestamp: "2025-09-15T12:52:00Z"  
      unit: Count/Second  
      value: 0.6670891564657616  
  region: us-west-2  
  timestamp: "2025-09-15T12:53:49.034106466Z"
```



```
input:  
  apiVersion: claude.fn.upbound.io/v1alpha1  
  kind: Prompt  
  systemPrompt: |  
    You are an intelligent RDS scaling system that analyzes XSQLInstance resources  
    and makes scaling decisions based on CloudWatch performance metrics.  
  userPrompt: |  
    You are analyzing an XSQLInstance resource for potential RDS scaling needs.
```

#### SCALING ANALYSIS CRITERIA:

1. Check performance metrics in status.performanceMetrics
2. Check current instanceClass in spec.parameters.instanceClass
3. Check if analysis was done recently (intelligent-scaling annotations)

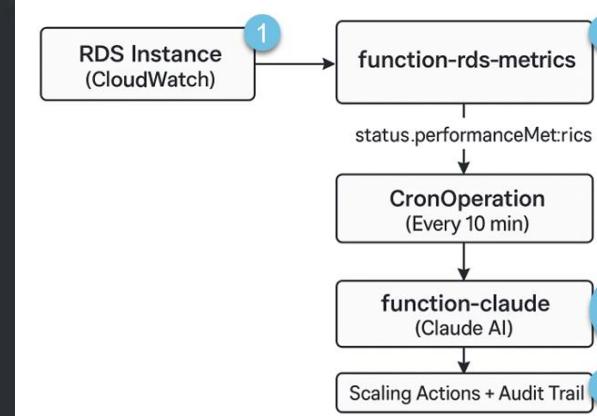
#### RATE LIMITING:

- Skip analysis if "intelligent-scaling/last-analyzed" annotation exists and is < 5 minutes old
- Only proceed if no recent analysis or if metrics show significant changes

#### SCALING TRIGGERS (conservative for cost control):

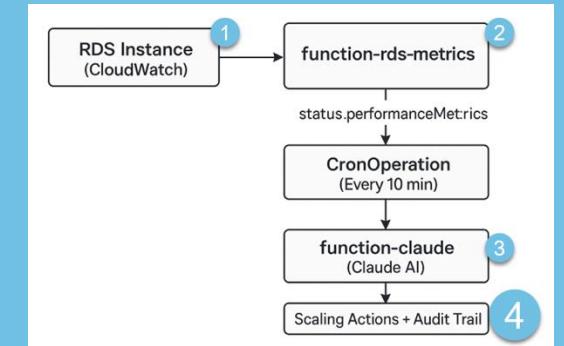
- SCALE UP: CPU > 85%, Memory < 15%, Connections > 85%
- SCALE DOWN: CPU < 20% AND Memory > 60% AND Connections < 30% for sustained period

#### INSTANCE CLASS PROGRESSION (bidirectional):



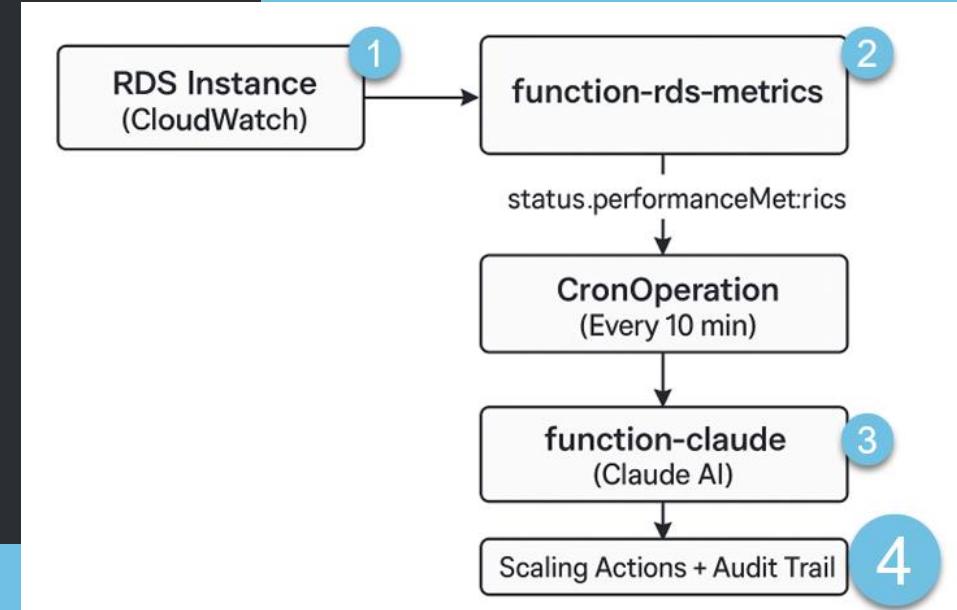
```

→ configuration-aws-database-ai git:(demo-sofia) ✘ k get xsqlinstances.aws.platform.upbound.io rds-metrics-database-ai
-scale -o yaml | yq
apiVersion: aws.platform.upbound.io/v1alpha1
kind: XSQLInstance
metadata:
  annotations:
    intelligent-scaling/last-analyzed: "2025-09-18T12:19:00Z"
    intelligent-scaling/last-scaled: "2025-09-17T23:59:59Z"
    intelligent-scaling/last-scaled-decision: No scaling needed. Very low resource utilization (CPU 2.09%, 0 connections) and already at smallest instance class (db.t3.micro).
    kubectl.kubernetes.io/last-applied-configuration: |
      {"apiVersion":"aws.platform.upbound.io/v1alpha1","kind":"XSQLInstance","metadata":{"annotations":{},"labels":{"scale":"me"},"name":"rds-metrics-database-ai-scale"},"spec":{"compositionSelector":{"matchLabels":{"type":"rds-metrics"}}, "parameters":{"autoGeneratePassword":true,"engine":"mariadb","engineVersion":"10.11","instanceClass":"db.t3.micro","networkRef":{"id":"rds-metrics-database-ai-scale"},"passwordSecretRef":{"key":"password","name":"mariadbsecret","namespace":"default"}, "publiclyAccessible":true,"region":"us-west-2","storageGB":5}, "writeConnectionSecretToRef":{"name":"rds-metrics-database-ai-scale","namespace":"default"}}}
```



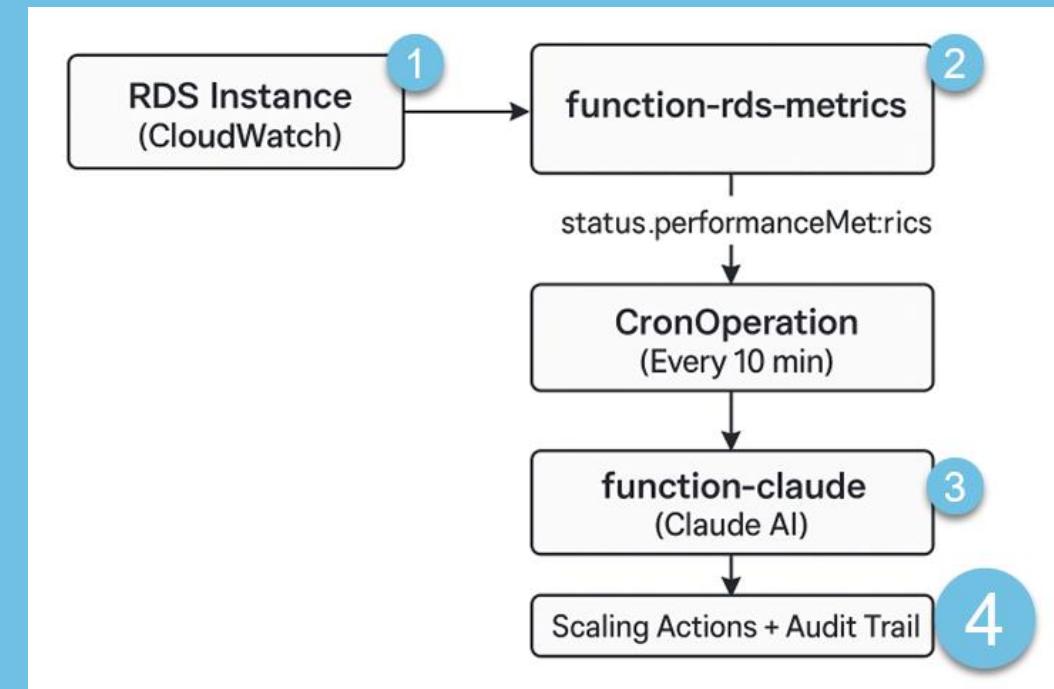
```
→ configuration-aws-database-ai git:(demo-sofia) ✘ ./perf-scale-demo.sh
📝 Starting DEMO load test (optimized for speed) ...
⌚ Load test running... Expected timeline:
- 30-60 seconds: CPU should hit 50%
- 1-2 minutes: CloudWatch metrics update
- 2-3 minutes: Claude analysis and scaling decision
- 5-10 minutes: Instance scaling completion

☰ Demo Check 1 (14:19:47) ☰
🔥 CPU: 1.8998733417772147% (threshold: 50%)
error: flag needs an argument: 'o' in -o
See 'kubectl get --help' for usage.
💻 Instance:
💻 Claude: ...
🎉 SCALING SUCCESSFUL! Instance upgraded to
🔴 Stopping load test...
✅ Demo complete!
```



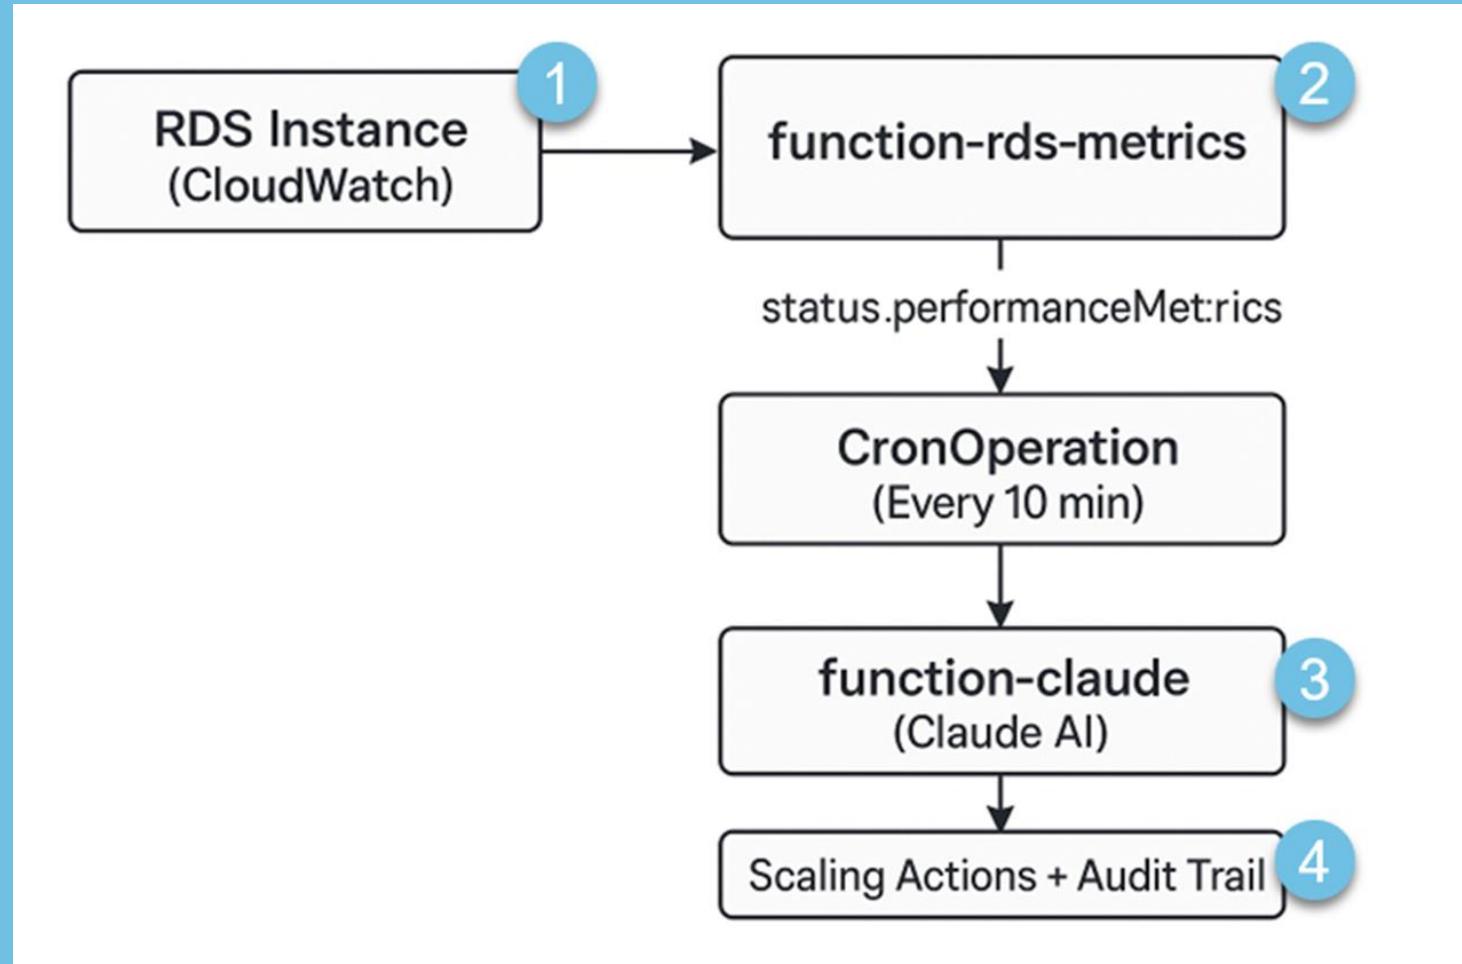
```
→ configuration-aws-database-ai git:(main) ✘ k get xsqlinstance.aws.platform.upbound.io/rds-metrics-database-ai-scale -o yaml|yq  
apiVersion: aws.platform.upbound.io/v1alpha1  
kind: XSQLInstance  
metadata:  
  annotations:  
    intelligent-scaling/last-analyzed: "2025-09-15T19:33:59Z"  
    intelligent-scaling/last-scaled: "2025-09-15T19:37:58Z"  
    intelligent-scaling/last-scaled-decision: Scaling up due to high CPU utilization (100%) and low freeable memory
```

```
parameters:  
  autoGeneratePassword: true  
  deletionPolicy: Delete  
  engine: mariadb  
  engineVersion: "10.11"  
  instanceClass: db.t3.medium  
  networkRef:  
    id: rds-metrics-database-ai-scale  
  passwordSecretRef:  
    key: password  
    name: mariadbsecret  
    namespace: default  
  providerConfigName: default  
  publiclyAccessible: true  
  region: us-west-2  
  storageGB: 5
```

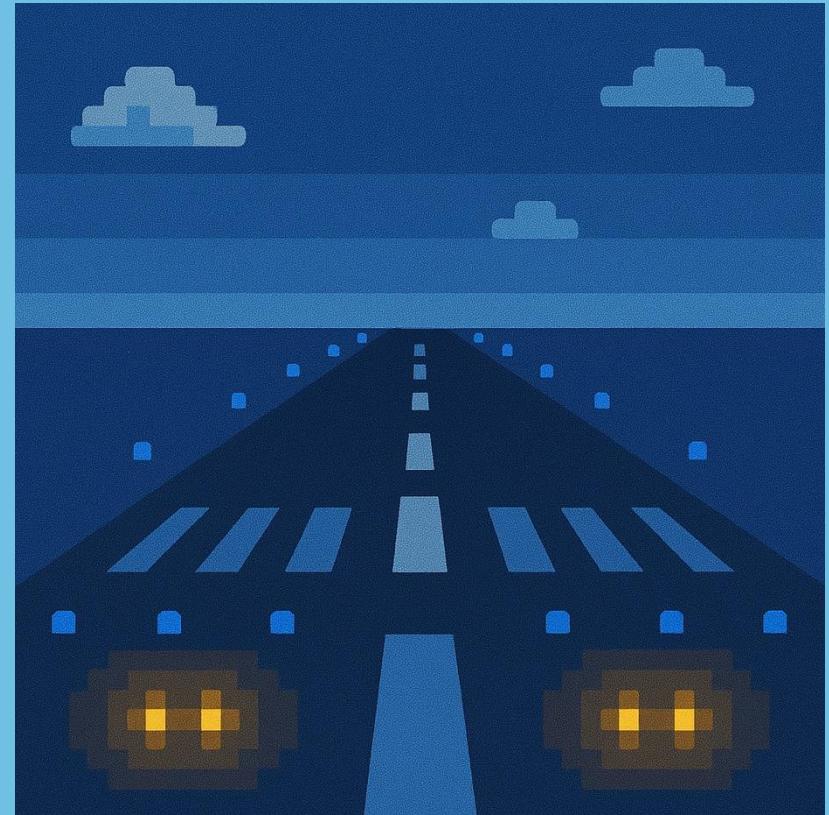


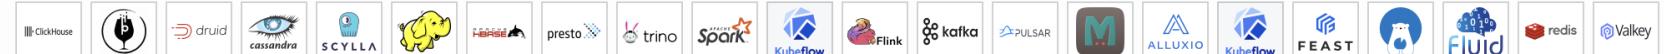
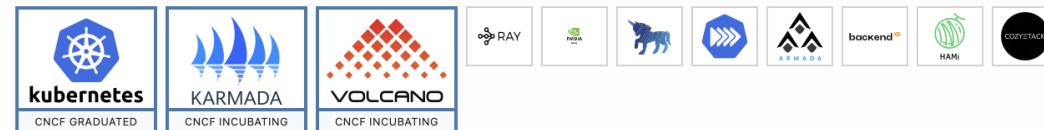
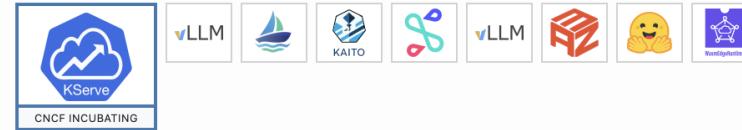
# Demo: An Autonomous Infrastructure Agent on Kubernetes

- A real AI agent:
  - Reads metrics
  - Analyzes them with Claude
  - Makes scaling decisions
  - Updates desired state
  - Crossplane enforces it



Landscape



Data Architecture CI/CD - Delivery General Orchestration Workload Observability Governance, Policy & Security AutoML ML Serving Vector Databases Distributed Training Agentic AI Data Science Open Enterprise AI Blue... Model/LLM Observability 

## Ray: Distributed Compute for Agent Workloads

- Ray is a distributed execution engine for Python-based agents
- Scales agent tasks, simulations, and tool-calling workloads
- Ideal for multi-step, long-running, or parallel reasoning flows
- Complements, not replaces, control-plane-native agents
- Works alongside Kubernetes and Crossplane in production stacks

## Where Open-Source Tool Fits in the Agent Stack

**Perception** → Prometheus, CloudWatch

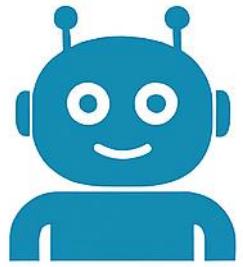
**Reasoning** → Claude, LangChain

**Coordination** → Autogen, CrewAI

**Distributed Execution** → Ray

**Action Layer** → Crossplane

**Safety / Guardrails** → Kyverno, Falco



### SRE Agent

Monitors system health and  
remediates incidents



### DevOps Provisioning Agent

Automates infrastructure  
provisioning and configuration

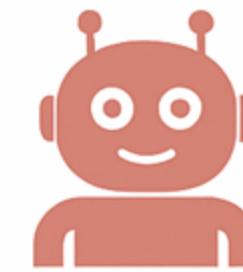


### Crossplane



### Security Compliance Agent

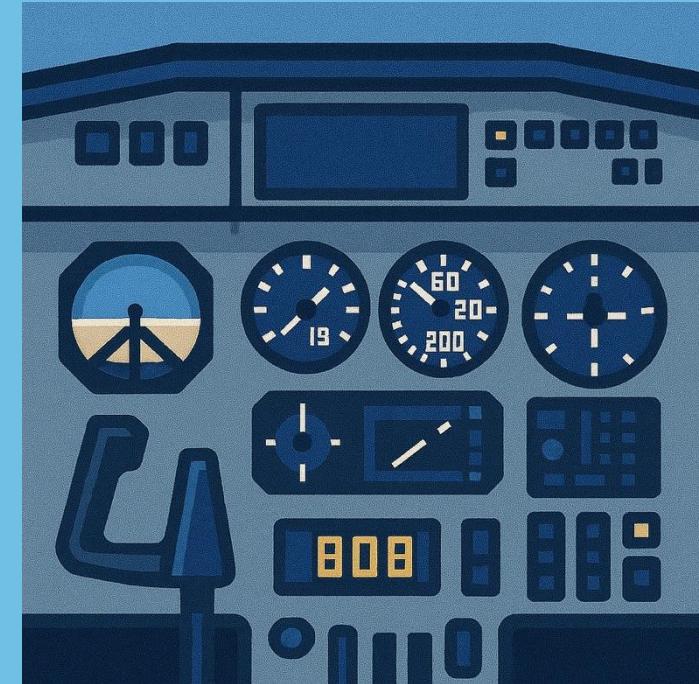
Enforces security policies and  
compliance requirements



### Cost Optimization Agent

Analyzes resource policies and  
compliance requirements

## Tips To Keep in Mind



## “It worked on my machine”, but now with AI

- **No Unified Observability:** Demos show console logs; production spans cloud services. We need metrics, tracing, and logging (Prometheus, OpenTelemetry) plus drift reconciliation (Crossplane).
- **No Human in the Loop:** Agents bypass reviews and playbooks. Guardrails and policies must enforce encryption, resource limits, and compliance automatically.
- **Speed & Scale of Incidents:** Agents act fast, so errors and compromises can spread in seconds.

# Policy and Guardrails for Autonomous Infrastructure

- **Preventive Guardrails (Admission Control & Policy Enforcement):** Kyverno policies act at deploy time: require resource limits, enforce trusted images, mandatory labels, disallow privileged containers. Policies-as-code in YAML provide automated compliance checks before changes go live.
- **Detective/Reactive Guardrails (Monitoring & Auto-Remediation):** Falco monitors runtime behavior: detects abnormal syscalls, privilege escalation, or suspicious network activity. Automated response via Falco Sidekick and Istio: isolate workloads, restart pods, and alert humans. Crossplane reconciliation adds drift correction for config tampering.

# Architectural Patterns for AI-Native Security

- Shift security left
- GitOps for the win
- Minimize agentic access
- Defense in depth

## Summary

- AI agents need structure, not scripts
- The control plane is the interface for AI
- Metrics (Prometheus) provide perception
- LLMs provide reasoning; Crossplane provides safe action
- Autonomous loops are possible today with open-source tools
- Policies & guardrails make autonomy safe

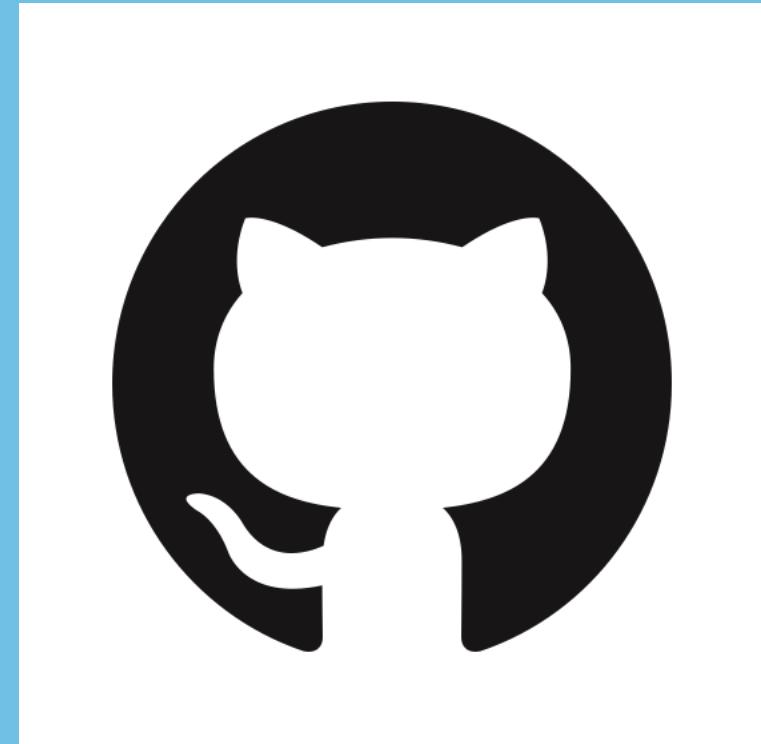
# Resources



[upbound.io/intelligent-control-plane](https://upbound.io/intelligent-control-plane)

```
performanceMetrics:  
  databaseName: rds-metrics-database-ai-scale  
  metrics:  
    CPUUtilization:  
      timestamp: "2025-09-15T12:52:00Z"  
      unit: Percent  
      value: 2.6191111704257306  
    DatabaseConnections:  
      timestamp: "2025-09-15T12:52:00Z"  
      unit: Count  
      value: 0  
    FreeStorageSpace:  
      timestamp: "2025-09-15T12:52:00Z"  
      unit: Bytes  
      value: 2775519232  
    FreeableMemory:  
      timestamp: "2025-09-15T12:52:00Z"  
      unit: Bytes  
      value: 188219392  
    ReadIOPS:  
      timestamp: "2025-09-15T12:52:00Z"  
      unit: Count/Second  
      value: 0.01667722891164404  
    WriteIOPS:  
      timestamp: "2025-09-15T12:52:00Z"  
      unit: Count/Second  
      value: 0.6670891564657616  
  region: us-west-2  
  timestamp: "2025-09-15T12:53:49.034106466Z"
```

[github.com/upbound/configuration-aws-database-ai](https://github.com/upbound/configuration-aws-database-ai)



[github.com/AnnieTalvasto/presentations](https://github.com/AnnieTalvasto/presentations)

Thank you!

