



Securing the Invisible Hands: Policy and Guardrails for AI-Driven Infrastructure



Annie Talvasto

- Sr Manager, Product Marketing & Community at Upbound
- CNCF Ambassador
- Azure MVP

- CloudGossip podcast & TechCraft Show
- Previously: CloudNativeLive & Program Chair for Secure AI Summit

Karl Ots

- Head of Cloud Security at EPAM
- Microsoft Security MVP

- Author of Securing Microsoft Azure OpenAI (Wiley)
- Cloud Gossip podcast

@AnnieTalvasto

@KarlOts.com

**How do we secure infrastructure that AI agents
create without human oversight?**

AI-Driven Infrastructure Risks

No Human Gatekeepers

AI makes rapid infra changes without manual review.

Misconfigurations or vulnerabilities can slip through.

Speed & Scale of Change

Autonomous changes happen too fast for human intervention if something goes wrong.

Propagating Impact


Misconfigurations or malicious actions by compromised agents can propagate widely before detection.

Need for guardrails

We require automated policies & monitors that work at machine speed to catch issues **in real-time**.


Policy as Code guardrails

- Policy as Code: Define security rules as declarative policies, preventing misconfigurations or non-compliant resources from ever being created
- Kubernetes Policies: e.g. Kyverno or OPA enforce rules on cluster changes before they take effect.
- Crossplane Composition Policies: extend policy scope from k8s cluster to cloud resources.




Crossplane 2.0: Declarative Control Plane for Safe Automation

- Crossplane: Open-source control plane (Kubernetes-based) to manage cloud infrastructure as code (YAML) across providers.
- **Composition Functions** (New in 2.0): Insert custom logic/policies into provisioning. (E.g., auto-attach mandatory security configs, or require approval for certain requests.)
- **Continuous Reconciliation**: Crossplane's controllers constantly compare declared state vs. actual state. If an AI-driven change causes drift (deviation), Crossplane detects & fixes it within seconds (auto-remediation).



Real-Time Monitoring and Response

- **Continuous Monitoring:** Even with good policies, runtime issues can occur. Use tools like Falco to watch running resources for suspicious behavior.
- **Automated Response:** Integrate alerts with actions: e.g., Istio/Network Policies to isolate a compromised service instantly.
- **Drift Detection:** Monitor config drift continuously. If any config deviates (e.g., AI agent changes something out-of-band), flag it or auto-roll back.



Architectural Patterns for AI-Native Security

Shift
security left

GitOps for
the win

Minimize
agenetic IAM

Defense in
Depth

Demo: Guardrails and Policy as Code in action



PROBLEMS

OUTPUT

DEBUG CONSOLE

TERMINAL

PORTS 5

```
• $ helm repo add kyverno https://kyverno.github.io/kyverno/
  helm repo update
  "kyverno" has been added to your repositories
  Hang tight while we grab the latest from your chart repositories...
  ...Successfully got an update from the "kyverno" chart repository
  Update Complete. *Happy Helming!*
  • $ helm install kyverno kyverno/kyverno --namespace kyverno --create-namespace
```

I

A screenshot of a terminal window displaying a YAML configuration file named `require-resources.yaml`. The file is located in a directory structure starting with `5_3 kyverno > ! require-resources.yaml`. The code itself defines a ClusterPolicy named `require-resources` with the following specifications:

```
apiVersion: kyverno.io/v1
kind: ClusterPolicy
metadata:
  name: require-resources
spec:
  validationFailureAction: enforce
  rules:
    - name: validate-resources
      match:
        resources:
          kinds:
            - Deployment
      validate:
        message: "Resource requests and limits are required"
        pattern:
          spec:
            template:
              spec:
                containers:
                  - name: "*"
                    resources:
                      limits:
                        memory: "?*"
                        cpu: "?*"
```

The screenshot shows a code editor interface with a dark theme. On the left is a vertical toolbar with icons for file operations like new, open, save, and search, as well as other developer tools. The main area displays a YAML configuration file:

```
! bad-deployment.yaml ! restrict-image-registries.yaml X
5_3 kyverno > ! restrict-image-registries.yaml

1 apiVersion: kyverno.io/v1
2 kind: ClusterPolicy
3 metadata:
4   name: restrict-image-registries
5 spec:
6   validationFailureAction: enforce
7   rules:
8     - name: validate-registries
9       match:
10         resources:
11           kinds:
12             - Pod
13       validate:
14         message: "Only images from approved registries are allowed"
15         pattern:
16           spec:
17             containers:
18               - name: "*"
19                 image: "nginx* | mcr.microsoft.com/*"
20
```

PROBLEMS OUTPUT DEBUG CONSOLE

TERMINAL

PORTS 5

ception' to true.

💡 Note: There is a trade-off when deciding which approach to take regarding Namespace usions. Please see the documentation at <https://kyverno.io/docs/installation/#security-perability> to understand the risks.

- \$ kubectl get pods -n kyverno

NAME	READY	STATUS	RESTARTS	AGE
kyverno-admission-controller-5bcbdff469-pnstk	0/1	Running	0	14s
kyverno-background-controller-7c7d4dbbc9-s6d9c	1/1	Running	0	14s
kyverno-cleanup-controller-745cbc6f8d-dcbh9	0/1	Running	0	14s
kyverno-reports-controller-7867ffd654-l9nkn	1/1	Running	0	14s

- \$ kubectl get pods -n kyverno -w

NAME	READY	STATUS	RESTARTS	AGE
kyverno-admission-controller-5bcbdff469-pnstk	1/1	Running	0	27s
kyverno-background-controller-7c7d4dbbc9-s6d9c	1/1	Running	0	27s
kyverno-cleanup-controller-745cbc6f8d-dcbh9	1/1	Running	0	27s
kyverno-reports-controller-7867ffd654-l9nkn	1/1	Running	0	27s

- ^C\$ code require-resources.yaml

- \$ kubectl apply -f require-resources.yaml

Warning: Validation failure actions enforce/audit are deprecated, use Enforce/Audit ins

.

clusterpolicy.kyverno.io/require-resources created

- \$ kubectl get clusterpolicy

NAME	ADMISSION	BACKGROUND	READY	AGE	MESSAGE
require-resources	true	true	True	15s	Ready

- \$

PROBLEMS 1 OUTPUT DEBUG CONSOLE

TERMINAL PORTS 5

```
• $ kubectl apply -f bad-deployment.yaml
Error from server: error when creating "bad-deployment.yaml": admission webhook "validate.knative.dev/v1beta1" denied the request:
resource Deployment/default/nginx-bad was blocked due to the following policies
require-resources:
  validate-resources: 'validation error: Resource requests and limits are required.
    rule validate-resources failed at path /spec/template/spec/containers/0/resources/limit
s/'
```

PROBLEMS 2 OUTPUT DEBUG CONSOLE TERMINAL PORTS 5

• \$ helm search repo kyverno-policies

NAME

CHART VERSION

APP VERSION

DESCRIPTION

kyverno/kyverno-policies

3.5.1

v1.15.1

Kubernetes Pod Securi

ndards implemented a...

• \$ helm install kyverno-policies --namespace kyverno kyverno/kyverno-policies --set pod
ityStandard=restricted --set validationFailureAction=Enforce

NAME: kyverno-policies

LAST DEPLOYED: Wed Sep 10 19:47:39 2025

NAMESPACE: kyverno

STATUS: deployed

REVISION: 1

TEST SUITE: None

NOTES:

Thank you for installing kyverno-policies 3.5.1 😊

We have installed the "restricted" profile of Pod Security Standards and set them in E mode.

Visit <https://kyverno.io/policies/> to find more sample policies.

◦ \$ █

PROBLEMS 2 OUTPUT DEBUG CONSOLE TERMINAL PORTS 5

We have installed the "restricted" profile of Pod Security Standards and set them in Enforce mode.

Visit <https://kyverno.io/policies/> to find more sample policies.

• \$ kubectl get clusterpolicies

NAME	ADMISSION	BACKGROUND	READY	AGE	MESSAGE
add-ns-labels	true	true	True	17m	Ready
disallow-capabilities	true	true	True	14s	Ready
disallow-capabilities-strict	true	true	True	14s	Ready
disallow-host-namespaces	true	true	True	14s	Ready
disallow-host-path	true	true	True	14s	Ready
disallow-host-ports	true	true	True	14s	Ready
disallow-host-process	true	true	True	14s	Ready
disallow-privilege-escalation	true	true	True	14s	Ready
disallow-privileged-containers	true	true	True	14s	Ready
disallow-proc-mount	true	true	True	14s	Ready
disallow-selinux	true	true	True	14s	Ready
require-resources	true	true	True	22m	Ready
require-run-as-non-root-user	true	true	True	14s	Ready
require-run-as-nonroot	true	true	True	14s	Ready
restrict-apparmor-profiles	true	true	True	14s	I Ready
restrict-seccomp	true	true	True	14s	I Ready
restrict-seccomp-strict	true	true	True	14s	Ready
restrict-sysctls	true	true	True	14s	Ready

! privileged-pod.yaml 1, U X

! root-pod.yaml 1, U

! secure-pod.yaml U

5_4 harden > ! privileged-pod.yaml > {} spec > [] containers > {} 0 > {} securityContext

io.k8s.api.core.v1.Pod (v1@pod.json)

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: privileged-pod
5 spec:
6   containers:
7     - name: nginx
8       image: nginx:1.19.8
9       securityContext:
10      privileged: true
11
```

PROBLEMS 2

OUTPUT

DEBUG CONSOLE

TERMINAL

PORTS 5

◦ \$ █

I

PROBLEMS 2 OUTPUT DEBUG CONSOLE TERMINAL PORTS 5

```
$ kubectl apply -f privileged-pod.yaml
Error from server: error when creating "privileged-pod.yaml": admission webhook "validate.k
yverno.svc-fail" denied the request:
resource Pod/default/privileged-pod was blocked due to the following policies
```

disallow-capabilities-strict:

require-drop-all: 'validation failure: Containers must drop 'ALL' capabilities.'

disallow-privilege-escalation:

privilege-escalation: 'validation error: Privilege escalation is disallowed. The fields spec.containers[*].securityContext.allowPrivilegeEscalation, spec.initContainers[*].securityContext.allowPrivilegeEscalation, and spec.ephemeralContainers[*].securityContext.allowPrivilegeEscalation must be set to 'false'. rule privilege-escalation failed at path /spec/containers/0/securityContext/allowPrivilegeEscalation/'

disallow-privileged-containers:

privileged-containers: 'validation error: Privileged mode is disallowed. The fields spec.containers[*].securityContext.privileged and spec.initContainers[*].securityContext.privileged must be unset or set to 'false'. rule privileged-containers failed at path /spec/containers/0/securityContext/privileged/'


require-run-as-nonroot:

run-as-non-root: 'validation error: Running as root is not allowed. Either the field spec.securityContext.runAsNonRoot must be set to 'true', or the fields spec.containers[

Demo: what did we see?

- AI agent submitted a Deployment without resource limit
- Kyverno validated this against the deployed ClusterPolicy
- Non-compliant Deployment was blocked

Bringing It All Together




Summary

- Autonomy != chaos
- Policy as Code for all
- Continuous compliance > periodic checks
- With the right design, we can enable AI confidently



upbound.io/intelligent-control-plane



a.co/d/ij4YA3V



Course

Microservices and API Security for Security Architects: From Gateway Protection to Container Security

By: Karl Ots • 1 month ago

linkedin.com/learning/microservices-and-api-security-for-security-architects-from-gateway-protection-to-container-security/

A pixelated illustration of a city skyline at sunset. The sky is a warm orange, with a large sun partially visible on the left. In the center, the Eiffel Tower is rendered in a dark brown color, appearing as a silhouette against the bright background. The city buildings are also in shades of brown and orange, creating a cohesive pixelated aesthetic.

Thank
you!